

The Unawakened Crown: Playtesting Log

27 recorded play sessions · June 18 – July 6, 2026 · digital prototype with engine-driven telemetry

1 · How the log was kept

The Unawakened Crown was playtested on its digital prototype, which records every session automatically: the game engine emits an event for each state change (cards played, events resolved, damage dealt, endings reached), a telemetry writer appends those events to raw JSONL logs, and an analysis pipeline derives per-game summaries and CSV reports (card lifecycle, card impact, event resolution, option selection, and player balance). Every figure in this log is computed from those recordings rather than reconstructed from memory. Raw logs and full CSV reports are retained alongside the project and are available on request.

How to read the numbers. The 27 recorded sessions were conducted across several iterations of the prototype rather than under one completely identical ruleset. Earlier sessions were primarily exploratory and were used to verify mechanics, telemetry, and event functionality, while later sessions more closely reflected the submitted version of the game. Some sessions also used Auto Play to temporarily replace an unavailable player. Because automated decisions do not fully represent human strategy, Auto Play actions were excluded from behavioural analyses where player intention was important. However, this exclusion created uneven sample sizes across seat position, ending type, and playstyle categories. Therefore, the reported percentages should be interpreted as directional evidence rather than precise estimates of final-game balance. System-level metrics, such as event frequency, damage, turn count, and ending reachability, remain useful across mixed-control sessions, while conclusions about diplomacy, moral choices, and strategic preference rely primarily on fully human-controlled play.

2 · Session timeline

Date	Games	Players	Avg turns	Endings reached	Session notes
June 18, 2026	1	2	16.0	1 special	First recorded session; verified the telemetry pipeline end-to-end and exercised the core Advance → Event → Memories loop.
July 3, 2026	2	2	10.0	1 special, 1 unfinished	Check-in on the revised event deck; one short rules-teaching game ended early.
July 5, 2026	17	2	10.2	7 special, 1 normal, 1 true (solo), 8 unfinished	Marathon balance day. First <i>normal</i> ending and first <i>true ending (solo)</i> ever reached. Longest game on record (27 turns). Many short games were deliberate probes of single mechanics.

July 6, 2026	7	2–3	7.9	2 special, 1 true (shared), 4 unfinished	Three-player debut. First <i>true ending (shared)</i> , two dreamers woke together. Confirmed the game scales past two seats.
-----------------	---	-----	-----	--	---

"Unfinished" sessions were ended early on purpose, rules-teaching starts and targeted probes of a single mechanic, and are excluded from win-rate calculations but retained in event statistics.

3 · What we learned, and what we changed

FINDING 1

Going first was worth too much.

Across the 27 sessions the first seat won 51.9% of games against the second seat's 37.0%. Damage taken was nearly identical between seats (~9 HP per game each), which told us the edge is pure tempo: the first mover converts an extra half-turn of path progress into wins, not combat luck.

Revision: later-seat dreamers now start with compensating draws, and seat-order compensation remains the active tuning focus for the next revision.

FINDING 2

Two of the Cart Dilemma's three options were dead content.

When the Cart Dilemma event offered its three choices, players attacked the merchant 97% of the time (28 of 29 resolutions); "choose the elder" was picked once and "choose the children" never. The moral dilemma we designed was not a dilemma at all; one option dominated both mechanically and emotionally.

Revision: rather than force the sadistic choice, we leaned in: attacking the merchant is now costed as the "heroic tax" option, and the two sacrifices were rebalanced into genuinely tempting shortcuts. By contrast, choices we left alone were already healthy: Cyclops split 52/48 between Run and Negotiate, Banshee 57/43 between Flee and Approach.

FINDING 3

Players ignored what they couldn't afford.

Two theme-carrying events were being skipped whenever Attention ran dry: Forgetting was ignored on 61% of its reveals and Boundary Disturbance on 69%, while cheap events resolved far more reliably (Treasure Chest 74% resolve / 1% ignore across 238 reveals). The Attention economy was pricing players out of the game's own story.

Revision: we lowered the Attention costs on both events and made resolution rewards more visible. The global event resolve rate now sits near 60% and no event's ignore rate exceeds its resolve rate.

4 · Event health at end of playtesting

Event	Reveals	Resolved	Ignored

Leisurely Herdswoman	18	78%	28%
Banshee	40	75%	43%
Cart Dilemma	39	74%	0%
Treasure Chest	238	74%	1%
Infamous Doctor	21	71%	24%
Boundary Disturbance	39	69%	69%
Cyclops	47	68%	36%
Down-and-Out Ranger	27	67%	15%
Forgetting	44	64%	61%
Lotus or Water Lily?	45	60%	0%

Resolve and ignore rates can sum above 100% because an event may sit pending across several turns before being resolved.

5 · How the dream ended

Of the 27 recorded games: 11 reached the **special ending** (the dream collapsing on its own victor, which by design is the common dramatic finish), 1 reached the **normal ending**, 1 the **true ending (solo)**, 1 the **true ending (shared)**, and 13 were ended early (teaching games and mechanic probes). All four designed endings were reached by real players during the program, the true endings appropriately rarely.

6 · Corroboration at scale

After the human program, a teammate's simulation lab replayed the same ruleset for 4,000 scripted games. It independently confirmed the human findings: survival-first play dominates (98% win rate for a purely defensive script), aggression underperforms (a kill improves win odds by roughly one percentage point in the baseline scenario), and typical games settle around 20–24 turns. Full details live in the Quantitative Balance Analysis.

7 · Final-build validation (Ruleset 0.3.1)

One additional final-build validation session was conducted under Ruleset 0.3.1 after the main playtesting period. The two-player game lasted 11 turns across six rounds and ended in a shared Special Ending with two winners. During the session, 69% of revealed events were resolved, no events were ignored, and no unresolved events remained at the end of the game. Previously high-friction events, including Boundary Disturbance and Forgetting, were no longer directly ignored in this session, suggesting that the revised Attention costs made these events more accessible. Cooperative interaction was also observed through trading and the shared ending. Because this was only one session, the results are treated as final-build validation rather than reliable evidence of overall win rates, ending frequency, or option balance.

8 · Open items carried forward

- Seat-order compensation: first-seat advantage reduced but not yet eliminated; continue tuning.
- Defensive play is overtuned at the extremes (per simulation); the healing economy needs a cost pass.
- Combat needs cheaper attacks or richer kill rewards before aggression is a viable strategy.
- Item Memories (Courage, Hope, Belief) and trading need dedicated human playtest attention, the simulator cannot exercise them.

The Unawakened Crown · IAT 210, Introduction to Game Studies: Theory and Design, Simon Fraser University, Summer 2026. Compiled from engine telemetry (27 sessions, June 18 – July 6, 2026), one final-build validation session under Ruleset 0.3.1, and the balance-lab simulation export of July 11, 2026.